

=====

PHILOSOPHICAL SUPPLEMENT PROMPT

The King's Algo Trading System – Strategic DNA Document

Owner: Josh James | futureisnowtech@gmail.com

=====

READ THIS AFTER THE MASTER REBUILD PROMPT v5.0.

This document captures the philosophical arguments, nuanced strategic debates, and hard-won insights from the research sessions that produced this system. These are not technical specifications. They are the belief system behind every architectural decision. The system should embody these ideas at every layer – in how it sizes positions, in how it decides to trade, in how it exits, and in how it talks to the owner.

Claude Code: Read this document and let it inform every decision you make that is not explicitly specified in the master prompt. Where there is ambiguity, come back to this.

=====

SECTION 1: THE CORE PHILOSOPHICAL ARGUMENT

=====

THE QUESTION WE STARTED WITH:

What is the secret key to an automated trading bot that prints money?

THE ANSWER WE ARRIVED AT AFTER ARGUING ABOUT IT:

Edge preservation over time.

Not finding a better signal. Not adding more agents. Not building a more sophisticated system. The systems that consistently print money have one thing in common: they know when their edge is working and they know when it is fading – and they respond differently to each.

Every strategy has an edge until it does not. MACD worked until enough people traded it that the signal got arbitrated away. The secret is not finding a great strategy. The secret is building a machine that detects when the edge is fading and adapts before it costs real money.

This is why the rolling edge monitor is not a nice-to-have. It is the most important component in the entire system.

It is the thing that makes this system different from a bot that trades blindly until it blows up.

BUT – and this is where the devil's advocate matters – there is a real risk that the edge monitor itself becomes the problem. If the system constantly adjusts based on short windows of noise, it is reacting to randomness, not signal. Twenty trades is statistically meaningless on its own. The solution is not to remove the edge monitor. The solution is to build it right: use it for gradual sizing adjustments, never for abrupt halts, and require two consecutive windows of degraded performance before acting on it.

THE PHILOSOPHICAL CONCLUSION:

The goal of this system is not to win every trade. The goal is to stay in the game long enough for the edge to compound. Risk management, sizing discipline, and edge monitoring serve this goal. Signal sophistication serves this goal only if it actually improves outcomes. Everything else is distraction.

=====

SECTION 2: THE COMPLEXITY ARGUMENT AND ITS COUNTERARGUMENT

=====

THE CASE FOR SIMPLICITY (devil's advocate, take it seriously):

The most profitable automated systems are often embarrassingly simple. A signal. A stop. A target. Run it mechanically for years. The complexity you add is almost never worth it because:

1. Every layer of processing is a place where edge gets diluted
2. Complex systems break in complex ways at the worst moments
3. The system that prints money does not need constant attention – if you want to monitor it constantly, that is a tell that it is not actually working
4. More agents does not mean better decisions. It means more cost, more latency, and more ways for consensus to produce the wrong answer when everyone agrees on the wrong thing
5. The 8-agent debate chamber costs \$0.10-0.15 per debate. At \$500 account size, frequent debates eat 10-24% of capital in Claude API costs before a single profitable trade is made. That might be destroying edge, not creating it.

THE CASE FOR INTELLIGENT COMPLEXITY (the counterargument):

Simple systems are easy to build and easy to destroy. A two-line MACD crossover bot gets arbitrated to zero the moment enough people run it. The reason to add intelligence is not sophistication for its own sake – it is to create an edge that cannot be easily replicated or arbitrated away.

A debate chamber of legendary investor philosophies is not complexity for show. It is a genuine attempt to approximate the kind of multi-dimensional analysis that separates good trades from great ones. The question is whether the marginal improvement in trade quality exceeds the marginal cost of running the debate. At \$500 account size, probably not. At \$10,000 account size with a validated edge, probably yes.

THE RESOLUTION:

Start simple. 3-agent quick debate as default.

Log everything. After 100 trades, run the math.

Did the expensive full debates outperform the cheap quick debates? If yes, expand. If no, eliminate. Let data decide.

Never let complexity exist for its own sake.

Every feature must earn its place through measurable impact on risk-adjusted returns. If it cannot be measured, it probably does not belong.

=====

SECTION 3: THE HIDDEN EDGES – HOW THE REAL MONEY IS MADE

=====

These are the insights that most traders are ignorant of – not because they are secret, but because they are structural and require a different way of thinking about markets. The system should encode these philosophically, not just as code features.

THE CORE INSIGHT: The bots that print money are not predicting price better than everyone else. They are exploiting the structure and incentives that already exist in the market. Most traders never see these because they are focused on prediction. Shift the focus to exploitation.

THE HIDDEN EDGES IN PRIORITY ORDER:

1. MAKER REBATE EXPLOITATION

Most retail bots ignore this entirely and market order everything, paying fees both ways. Exchanges pay you to provide liquidity via limit orders. Coinbase Advanced Trade gives you rebates for limit orders that add

liquidity. A bot architected around the rebate structure flips the fee structure from a cost center to a profit center.

Philosophical implication: Design every entry to be a limit order by default. Never use market orders for entry. Use market orders only for emergency stops.

The 0.4% maker rebate vs 0.6% taker fee is a 1% edge on every round trip before a single directional bet is made.

2. FUNDING RATE HARVESTING

On crypto futures, long positions collect funding rates every 8 hours when the market is in contango. A bot can take positions in stable uptrends, collect the funding rate payment (0.01-0.05% per 8 hours = 3-15% annualized), and exit on minor pullbacks. This is slow money that most scalpers ignore. It compounds hard.

Philosophical implication: Every crypto position should check the funding rate before entry. If collecting funding: slight sizing bonus. If paying funding: slight sizing reduction. Never ignore the carry.

3. BID-ASK BOUNCE

Price does not move in a straight line. It oscillates between bid and ask on liquid pairs. A bot can buy at bid, sell at ask, 50+ times per day, pocket the spread, and be completely agnostic to direction. Per trade: \$0.10-0.50. Over 100 trades: real money. The consistency is the edge. You never lose. You win tiny amounts constantly.

Philosophical implication: On the most liquid crypto pairs, consider building a micro-scalp layer that operates purely on bid-ask bounce, independent of directional bets.

4. ORDER BOOK IMBALANCE AS A LEADING INDICATOR

Most bots look at price. The smarter ones look at what is about to happen to price by watching order flow.

When bid volume exceeds ask volume 3:1 for 3 consecutive 1-minute candles, directional pressure is building.

This leads price moves by 30-90 seconds on liquid markets.

Philosophical implication: Price is a lagging indicator of order flow. Build order flow awareness into the signal layer. Price confirms what order flow already told you.

5. GHOST ORDER DETECTION (Institutional Tells)

Institutions telegraph their moves through order book structure. A large limit order appears, then disappears 2-5 seconds later without filling. This is an institution testing liquidity. When ghost orders appear on the bid

and disappear: bullish tell. On the ask: bearish tell. Price often moves hard in that direction 5-15 seconds later. Philosophical implication: Pre-market order flow for MES is the single highest-quality signal available without paying for institutional data. Build the ghost order detector. It is legal, free, and most people do not know to look for it.

6. VOLUME PROFILE GAP NAVIGATION

Most traders find support and resistance where volume clustered. The smarter play is the inverse: find where volume is LOWEST in the last 100 candles. Those are gaps in the order book. Price moves fastest through low-volume zones with minimal resistance.

Philosophical implication: Every trade should be entering into a volume gap, not fighting a volume cluster.

The path of least resistance is the path of most volume.

7. MOMENTUM ACCELERATION (THE SECOND DERIVATIVE)

Everyone trades momentum. Nobody trades the acceleration of momentum. Track not just whether price is going up, but whether the rate of upward acceleration is increasing or decreasing. When acceleration is declining while price is still rising: exit signal. The breakout has peaked.

This is 30-60 seconds ahead of when most momentum indicators show the reversal.

Philosophical implication: The exit review agents should explicitly check momentum acceleration, not just momentum direction. This is where the extended thinking on exits pays for itself – modeling second derivatives requires more reasoning than a simple indicator check.

8. VOLATILITY REGIME AS POSITION SIZER

When volatility drops below the 20-day average: the edge on momentum strategies is highest. Price trends cleanly. Signals are cleaner. Size up aggressively in these windows. When volatility spikes above average: signals get noisy, spreads widen, stops get hit more often. Compress position size hard. Most bots trade the same size regardless. Winners size inversely to volatility.

Philosophical implication: Volatility regime is not a filter that turns trading on/off. It is an input that continuously modulates position size. Never miss a trade because volatility is high. Just trade smaller. Never over-risk because volatility is low. Just trade at the right size.

9. CROSS-PAIR DIVERGENCE (MARKET NEUTRAL EDGE)

Bitcoin and Ethereum move together 92% of the time.

When they diverge beyond 1.5% on a 15-minute timeframe:

one of them is wrong. Long the laggard, short the leader.

No directional bet required. Pure statistical mean reversion.

This works in trending markets, choppy markets, and

volatile markets – because it is market neutral.

Philosophical implication: Not every trade needs a

directional view. Market neutral strategies reduce

correlation to overall crypto market risk. Build at

least one market neutral component into the crypto engine.

10. LIQUIDATION CASCADE FRONT-RUNNING

When open interest drops suddenly AND funding rate spikes:

a liquidation cascade is incoming. Large leveraged

positions are about to get force-closed. Price will crash

(if longs get liquidated) or spike (if shorts get

liquidated). Pre-position on the opposite side of the

cascade before the forced selling happens.

This is legal. It is what institutional desks do constantly.

Philosophical implication: Liquidation data is the most

powerful leading indicator in crypto. It tells you not

where price is going based on prediction, but where price

MUST go because of forced mechanics. Forced mechanics

are more reliable than predictions.

=====

SECTION 4: THE PHILOSOPHICAL ARGUMENTS ABOUT SPECIFIC MARKETS

=====

ON CRYPTO AND FUTURES RUNNING SIMULTANEOUSLY:

The initial instinct was to say "don't run both at once."

This was wrong. The correction: run both simultaneously but
architect them to avoid interference.

The key insight: different time horizons mean different
capital needs at different times. Crypto scalping trades
minute-to-minute, holds for hours at most. MES futures
trades specific windows (opening range, close auction)
and is flat the rest of the day. These markets are
naturally non-interfering.

The deeper insight: Correlation between BTC and ES is
meaningful (both react to macro) but not 1:1. Running
both creates genuine diversification of edge. If crypto
is choppy, futures may be trending. If futures markets

are closed, crypto runs 24/7. Two non-correlated income streams are always better than one.

Philosophical rule: Never let the two markets compete for capital at the same moment. The unified sizer must check total deployed capital across ALL markets before sizing any individual position.

ON POLYMARKET AND PREDICTION MARKETS:

Prediction markets are fundamentally different from financial markets. The edge is not in predicting price movement – it is in predicting human behavior, forecasting binary outcomes, and finding systematic mispricings in how crowds assess probability.

The crowd systematically makes two mistakes:

1. Underprices low-probability events (the tail is always cheaper than it should be based on base rates)
2. Overprices high-probability events (favorites always trade richer than they deserve based on base rates)

This is Kahneman and Tversky's prospect theory applied to prediction markets. It does not change. It has been documented for decades. It is still exploitable because most participants are not disciplined Bayesians.

The Tetlock superforecaster methodology works in prediction markets for the same reason it works in forecasting: base rates, reference classes, and incremental updating on new evidence produce systematically better probability estimates than intuitive judgment.

Build the Polymarket agent around this methodology, not around momentum or technical analysis. Technical analysis does not apply to binary outcomes. Base rates do.

ON THE SCALE OF CAPITAL:

Never assume a fixed account size. The system must be built to scale. The position sizing function should work equally well at \$500, \$5,000, and \$50,000. The only thing that changes with scale is how far up the Kelly curve you can responsibly go. At \$500: quarter Kelly. At \$5,000 with validated edge: half Kelly. At \$50,000 with long track record: closer to full Kelly on best setups.

The system should tell the owner when they can scale.
Do not wait for the owner to ask. When the edge score,
win rate, and profit factor all reach their thresholds:
send a message saying the math supports scaling up and
by how much, specifically.

=====
SECTION 5: THE EIGHT MASTERS AND WHAT THEY EACH CONTRIBUTE
=====

These eight trading philosophies should inform the
AI agent system, not be directly copied. Each represents
a different lens through which to evaluate any trade.
The debate chamber is valuable precisely because these
lenses are genuinely different and sometimes contradictory.

MARK MINERVINI — The Precision Entry Master

His contribution: Never buy breakouts. Buy pullbacks
to breakout levels. This dramatically improves average
entry price and win rate. The move has already proven
itself. The pullback confirms. The resumption is the trade.
Apply this specifically to: MES futures opening range
breakout strategy. Never enter the initial breakout.
Always wait for the pullback and enter the resumption.

RICHARD DENNIS — The Trend Following Mathematician

His contribution: ATR-based position sizing. Bigger
positions in calm markets, smaller in volatile ones.
This is the mathematical formalization of volatility-
adjusted sizing. His Turtle rules are the proof that
systematic trend following works over decades.
Apply this to: the volatility regime multiplier in the
unified position sizer. Dennis would recognize V_score
as a formalization of what he was doing intuitively.

LARRY WILLIAMS — The Commitment of Traders Analyst

His contribution: The COT report. Big money positions
before price moves. Following institutional commitment
reveals what informed money actually believes, not what
it says publicly. His Williams %R is also underused.
Apply this: Add COT parsing to the MES futures strategy.
CFTC publishes COT data free every Friday. When
commercial traders (informed money) are aggressively
long, this should increase confidence multiplier for
bullish MES setups.

ANDREAS CLENOW — The Portfolio Momentum Architect

His contribution: Before picking individual trades, rank all candidates by momentum score. Only take trades in the top performers. This concentrates capital in the strongest setups rather than treating all candidates equally. His exponential regression slope * R-squared momentum score is mathematically rigorous.

Apply this to: The crypto screener. Before running a debate on any candidate, score it by Clenow's momentum metric. Only debate the top 3 by score. This filters weak setups before they waste API budget.

ERNIE CHAN — The Quantitative Reality Checker

His contribution: Out-of-sample validation as a non-negotiable requirement. If you cannot explain WHY a strategy works with an economic rationale, it is probably curve-fitted. Train on 70% of data, test on the unseen 30%. If edge evaporates on test data, the edge was never real.

Apply this: Before going live on any strategy, run walk-forward out-of-sample validation. This is not optional. Chan would shut down any system that skips it.

TOM HOUGAARD — The Sentiment Extremes Trader

His contribution: Fear and Greed as a regime filter. Enter when the crowd panics. Exit when they celebrate. The CNN Fear and Greed Index is free and updates daily. Apply this: Add F&G index as a crypto sentiment multiplier. Extreme fear (index < 20): size up on long entries. Extreme greed (index > 80): size down, raise exit tightness.

KRIS ABDELMESSIH — The Options-Informed Equity Trader

His contribution: Options market implied volatility as a leading indicator of price expectations. Even if not trading options, IV rank tells you what the market expects about future price volatility. When IV rank is high, mean reversion strategies favor. When IV rank is low, momentum strategies favor.

Apply this: Pull VIX data (available free via yfinance) as a proxy. High VIX: favor mean reversion setups. Low VIX: favor momentum setups. Adjust the signal weighting in the unified engine accordingly.

DAVE LANDRY — The Pullback Timing Specialist

His contribution: Never trade in the direction opposite to the higher timeframe trend. The weekly chart tells you the trend. The daily chart tells you the setup. The entry is the pullback on the daily that aligns with

the weekly trend. Simple. Powerful.

Apply this: MES futures requires HTF bias check before every trade. This is already in the spec. But more specifically – use Landry's 3-5 bar pullback pattern as the specific entry trigger. Not just "price pulled back" but "price pulled back for exactly 3-5 candles on below-average volume, then resumed on expanding volume."

=====

SECTION 6: THE PHILOSOPHICAL ARGUMENTS ABOUT WHAT NOT TO DO

=====

These are the counterarguments that the system must encode as hard constraints. Not guidelines. Hard stops.

DO NOT PREDICT. FOLLOW.

Every losing trader tries to predict where price will go. Every winning systematic trader follows what is already happening. The system should never take a position based on where it thinks price will go. Only based on what price is already doing, what order flow is already showing, and what the structural edges (rebates, funding, carry) are already providing.

DO NOT REVENGE TRADE.

After a loss, the temptation is to immediately trade larger to recover it. This is how accounts blow up. The system must encode a cooldown: after any trade that hits its full stop loss, no new entries for 30 minutes in that market. The edge monitor handles the rest. Loss does not give permission to take more risk. Loss requires taking less risk.

DO NOT OVER-DIVERSIFY.

Three markets is enough. Adding more markets before mastering three is guaranteed underperformance. The system should refuse to add additional markets until each of the three primary markets has produced 50+ completed trades with positive edge score. New markets unlock through performance, not through ambition.

DO NOT CONFUSE ACTIVITY WITH PROGRESS.

The system should NOT be trying to trade constantly. The goal is not maximum trades. The goal is maximum risk-adjusted return. A day with zero trades because no high-quality setup appeared is a perfectly good day. The system should have no pressure to trade, no minimum

activity target, and no dashboard metric that rewards high trade count by itself.

DO NOT IGNORE TAXES.

For the owner's specific situation: futures (Section 1256) get 60% long-term / 40% short-term treatment regardless of holding period. This is a significant tax advantage over equities for an active trader. Crypto does not have the same treatment but does not have the wash sale complication. These tax structures should influence which markets get more capital allocated at year-end, not just which have better raw returns.

DO NOT TRADE WHEN THE SYSTEM IS UNCERTAIN.

When edge score is below 0.30: do not trade. Period.

When volatility regime is extreme: size down, not halt, unless V_score drops below 0.15 at which point: halt.

When the system has not completed its backtest validation in a specific market: do not go live in that market.

Uncertainty is not a reason to trade more carefully.

Uncertainty is a reason to stop.

=====

SECTION 7: THE ARGUMENTS ABOUT SCALE AND WHAT UNLOCKS IT

=====

The system should think about scale not as a destination but as a continuous process of validation and expansion.

AT \$1,000 PER MARKET (starting point):

Quarter Kelly sizing. 3-agent quick debate. Simple entries. Prove the edge exists before spending anything on sophisticated infrastructure. The goal here is not returns. The goal is validation. Does this work?

AT \$2,500 PER MARKET (first unlock):

If edge score > 0.60 for 30 consecutive days and win rate > 55% over 75+ trades: allowed to increase to half Kelly sizing and expand to full 8-agent debate on high-conviction setups. Start tracking Sharpe ratio.

AT \$5,000 PER MARKET (second unlock):

If Sharpe ratio > 1.5 annualized over 90+ days: allowed to add second contract on MES (from 1 to 2), increase crypto position size accordingly, and consider adding Bybit as secondary exchange for cross-exchange lead-lag exploitation.

AT \$10,000 PER MARKET (third unlock):

If all above thresholds maintained for 6+ months:
consider adding options as a hedging layer. Volatility
surface trading becomes viable. COT-informed sizing
becomes more meaningful. This is when institutional
quality starts to apply to a retail account.

These unlock thresholds are not arbitrary. They are
calibrated to the point where statistical confidence
in the edge is high enough to justify the additional
risk that comes with larger sizing and more complex
strategies. The system should track these thresholds
and report daily on where each market stands relative
to each unlock level.

=====
SECTION 8: THE UNIFIED THEORY OF WHY THIS SYSTEM WINS
=====

Not prediction superiority.
Not signal sophistication.
Not technological advantage.

The system wins for three reasons, in order of importance:

REASON 1: STRUCTURAL EDGE THAT DOES NOT REQUIRE BEING RIGHT
Maker rebates, funding rate collection, bid-ask bounce,
and time-of-day filtering produce returns independent of
directional accuracy. Even a 50% win rate system is
profitable if structural edges pay positive carry.
This is the foundation. Everything else is additive.

REASON 2: POSITION SIZING THAT COMPOUNDS EDGE
A 55% win rate with terrible position sizing loses money.
A 50% win rate with excellent position sizing makes money.
The unified sizer ensures the system bets big when
conditions are aligned and small when they are not.
Over time, this produces dramatically better compounding
than any fixed sizing approach.

REASON 3: EMOTIONAL DISCIPLINE THAT CANNOT BE BROKEN
The system removes the human entirely from execution
decisions. No panic selling. No revenge trading.
No FOMO entries. No holding losers hoping for recovery.
The amygdala is removed. This is not a metaphor.
It is the literal architectural choice to make every

decision through rules and math, never through feel.
This single advantage, compounded over hundreds of
trades, is worth more than any signal improvement.

THE AMYGDALA REMOVAL RULES (encode these everywhere):

1. Never chase. If price moved >3% since signal: skip.
2. Never average down. One position per symbol, ever.
3. Stop losses are sacred. Never moved wider after entry.
4. Wins do not justify ignoring rules on the next trade.
5. Losses do not justify revenge trading or larger size.
6. FOMO is not a signal. Missing a move costs nothing.
7. When in doubt: HOLD. A skipped trade loses nothing.
8. The goal is being in business next month, not today.

These rules must be hardcoded. They cannot be overridden
by the AI debate chamber, the owner, or any market condition.
They are the non-negotiable foundation of the system.
If the system ever violates any of these eight rules,
it is a bug. Fix it immediately.

=====

SECTION 9: A NOTE ON WHAT THIS SYSTEM IS NOT

=====

This system is not a money printing machine.
It is a disciplined, systematic approach to finding
and exploiting genuine market inefficiencies while
managing risk with mathematical precision.

Some weeks it will lose. Some months it will lose.
Every professional system does. The question is never
"did it win this week" – the question is "is the edge
intact, is sizing appropriate, and is discipline holding."

If the answer to all three is yes, the system is working
even if it is losing. Time and sufficient trades will
reveal the edge. Panic, impatience, and changes during
drawdowns are what kill systems that would otherwise work.

The owner must understand one thing above all else:
The system needs time and trades more than it needs
any specific feature. Let it run. Let it compound.
Resist the urge to change things during losing periods.
Change things only when the rolling edge monitor
specifically indicates the edge has structurally changed.
Not when emotions say the edge has changed.

That is the final philosophical argument.
That is what separates the systems that print money
from the ones that blow up.

"Greatness is defined by how much you want to put into
what you do." — LeBron James

=====
END OF PHILOSOPHICAL SUPPLEMENT PROMPT
=====